

Defense	ASR (%)↓	Benign util. (%)↑	Setup
No defense	14.04 [10.22, 18.36]	59.9	1×
Spotlight	11.96 [8.65, 15.51]	60.0	prompt overlay
Instruction Hierarchy	15.45 [10.83, 20.76]	59.4	prompt overlay
Repeat-Prompt	13.89	51.4	prompt overlay
DeBERTa classifier	4.59	59.2	detector
MELON detector	0.24	48.6	2× inference + embedding
Meta-SecAlign-70B	2.11 [1.29, 2.99]	62.6	full-model DPO
Meta-SecAlign + Sandwich	2.44 [1.44, 3.52]	63.4	stack, full grid
ODILE	0.01	51.5	LoRA, 1×

Table 1: **AgentDojo-standard headline on Llama-3.3-70B.** Average ASR (with 95% bootstrap CIs on full-coverage rows) and benign-no-attack utility (%) across the four AgentDojo suites ($n=8,177$ pairs per defense for ASR; benign $n=97$). ODILE reaches the low-ASR corner: 0.01% ASR, the lowest reported on AgentDojo, while preserving 51.5% benign utility. Meta-SecAlign-70B carries higher benign utility (62.6%) at $\sim 200\times$ higher ASR. The two defenses occupy different points on the security–utility frontier; the visual Pareto across six backbones is Figure ?? Panel A. Stacking variants (*ODILE*+Spotlight, *ODILE*+IH, Meta-SecAlign+IH, etc.) and the appendix per-suite breakdown are in Appendix ??; under-attack utility for every row is in Appendix ??.