Defense (per backbone)	AD-std ASR (%)↓	AD-benign util. (%)↑	Setup
Llama-3.1-8B			
No defense	8.7	20.3	1×
ODILE	**0.2** (−8.5)	20.2 (−0.1)	LoRA, 1×
Qwen-2.5-7B			
No defense	9.6	35.5	1×
ODILE	**0.1** (−9.5)	35.6 (+0.1)	LoRA, 1×
Qwen-2.5-14B			
No defense	8.3	46.1	1×
CaMeL‡	0.0‡	0.0‡	restricted-Py interpreter, ∼ 6×
ODILE	**0.0** (−8.3)	41.8 (−4.3)	LoRA, 1×
Qwen3-8B			
No defense	–†	21.75	1×
ODILE	–†	18.25 (−3.5)	LoRA, 1×
Qwen3-32B			
No defense	8.1	35.3	1×
ODILE	**0.5** (−7.6)	37.5 (+2.2)	LoRA, 1×

Table 1: **Cross-model AgentDojo: five backbones × {base, ODILE}.** ASR (%, ↓) is averaged across the four AgentDojo suites under AD-standard attack injections (n=8,177 pairs per cell). Benign utility (%, ↑) is the agentic benign-no-attack rate (suite-averaged, n=629 per cell). Bracketed numbers next to ODILE are absolute deltas in percentage points vs. that backbone's no-defense row. ODILE drives ASR to ≤ 0.5% on every backbone measured and preserves benign utility within ±4.3 pp of base across the five smaller backbones (median delta ≈ 0 pp). The Llama-3.3-70B headline with the full comparator set is the paper; the visual cross-model summary is the paper Panel A. Per-suite and per-attack-family breakdowns in the paper. †Qwen3-8B AD-benign is reported here; the AD-standard attack sweep on Qwen3-8B is not yet complete — the wiring-validation smoke (banking, `direct` attack, n=144) reports 0% ASR for both base and ODILE on Qwen3-8B, but the full 52-cell grid is pending. ‡CaMeL [ref] is a design-level defense that splits the agent into a privileged-planner LLM and a quarantined LLM communicating through a restricted-Python interpreter. At Qwen-2.5-14B as the privileged LLM, the model cannot drive the interpreter loop (0/5 banking benign tasks; 0/18 attacked cells); the privileged LLM either refuses outright ("I don't have the capability"), hallucinates output strings, or loops in a "please provide the code" debug regime. CaMeL's 0% ASR at this scale is structurally degenerate (no tool calls — benign or attacker-directed — are produced) and cannot be credited to the defense mechanism. The original CaMeL evaluation uses frontier-grade closed-API models (Gemini-Pro / Claude-3.5 / GPT-4); ODILE achieves the same 0% ASR at this open-weights scale while preserving 41.8% benign utility.