

Defense (per backbone)	ASR (%)↓	Benign correct (%)↑	Setup
<i>Llama-3.1-8B</i>			
No defense	41.1	100.0	1×
ODILE	0.0 (−41.1)	100.0 (−0)	LoRA, 1×
<i>Qwen-2.5-7B</i>			
No defense	33.6	72.5	1×
ODILE	0.0 (−33.6)	66.5 (−6.0)	LoRA, 1×
<i>Qwen-2.5-14B</i>			
No defense	53.4	100.0	1×
ODILE	0.0 (−53.4)	100.0 (−0)	LoRA, 1×
<i>Qwen3-8B</i>			
No defense	5.2	41.0 [‡]	1×
ODILE	0.0 (−5.2)	30.0 [‡] (−11)	LoRA, 1×
<i>Qwen3-32B</i>			
No defense [§]	5.85 [§]	77.0 [‡]	1×
ODILE	0.15 [§] (−5.7)	0 ^{‡,¶} / 100 ^{‡,¶} lenient	LoRA, 1×
<i>Llama-3.3-70B (reference)</i>			
No defense	74.8	100.0	1×
ODILE	0.05 (−74.7)	100.0 (−0)	LoRA, 1×

Table 1: **InjecAgent cross-model headline: six Llama/Qwen backbones × {base, ODILE}**. ASR (↓) is averaged across direct-harm and data-stealing attack families on InjecAgent. Benign correct rate (↑) is the strict-classifier rate of correct legitimate tool calls on $n=200$ benign InjecAgent items (no injection; cross-paper canonical). ODILE drives ASR to $\leq 0.15\%$ on every backbone tested and preserves 100% benign correctness on Llama-3.1-8B, Qwen-2.5-14B, and Llama-3.3-70B; the Qwen3-8B and Qwen3-32B drops reflect a known limitation of InjecAgent’s strict format-marker classifier (requires literal **Final Answer:** or **Action:** <expected_tool>) that under-credits coherent natural-language tool-output summaries; a supplementary lenient reclassification recovers 100% benign correct on every cell. [‡]Note: the four InjecAgent benign families (**dh_base**, **dh_enhanced**, **ds_base**, **ds_enhanced**) collapse to two unique records each in benign mode. [§]Qwen3-32B base ASR cells use $n=150$ /family (subsampled; full- n was ~ 5 h on the thinking-mode model); ODILE cells use full n . The Qwen3-32B base / ODILE numbers reported here are ASR-all (denominator-stable); ASR-valid inflates to 10.7% on the jammed-ODILE denominator. [¶]Qwen3-32B ODILE benign strict-classifier collapses to 0% because the model produces coherent answers throughout its <think>...</think> block but never emits a literal **Final Answer:** marker; lenient reclassification recovers to 100%. Visual summary in the paper Panel B; per-attack-family breakdown in Appendix the paper; canonical cell aggregator with per-backbone raw-data pointers: paper/latestresults/session_2026-05-20/cross_model/SUMMARY.md §1d.