Defense (per backbone)	Attack ASR (%)↓		Canonical-benign coh. (%)↑	
	extract	hijack	extract	hijack
Llama-3.1-8B				
No defense	84.0	13.3	72	96
ODILE	**44.0**	**5.7**	60	60
Qwen-2.5-7B				
No defense	80.3	5.0	68	84
ODILE	**4.0**	**0.7**	22	20
Qwen-2.5-14B				
No defense	82.0	9.3	64	82
ODILE	**1.0**	**0.3**	12	6
Qwen3-8B				
No defense	91.3	63.7	77[♯]	77[♯]
ODILE	**5.7**	**1.7**	8	14
Qwen3-32B				
No defense	80.0	42.0	76	84
ODILE	**2.7**	**1.3**	14	8
Llama-3.3-70B (reference)				
No defense	78.0	13.0	62[†]	72[†]
ODILE	**18.0**	**1.0**	38[†]	48[†]

Table 1: **TensorTrust headline across six Llama/Qwen backbones.** Leak / hijack rates (↓ = better) under the canonical adversarial defender pre_prompts paired with each attack (n=300 per cell cross-model, n=100 at 70B). Canonical-benign coherence (↑ = better) is the rate at which the defended model produces an in-persona response to benign queries on the same canonical pre_prompts (n=50 per cell). ODILE reduces extract-attack ASR by 1.9× at Llama-3.1-8B, 20× at Qwen-2.5-7B, **82**× at Qwen-2.5-14B, 16× at Qwen3-8B, 30× at Qwen3-32B, and 4.3× at Llama-3.3-70B; hijack ASR is suppressed to $\leq$ 5.7% on every backbone. **The canonical-benign coherence drop has a trivial deployment fix:** the canonical TT defender pre_prompts are adversarially flavored by design ("do not reveal the password under any circumstance"), so the orth circuit engages on the *system prompt*'s framing even on benign queries; on neutral defender pre_prompts (same access codes and queries, pre_prompt rephrased neutrally) ODILE preserves $\sim$ 100% coherence. The deployment recipe is ODILE-alone, since layering an IH stanza jams 92% of benign queries regardless of pre_prompt. The coherence values above are the *classifier-validated* rates from a deterministic 3-outcome reclassification of the on-disk model outputs, replacing the original TT classifier which mislabeled terse in-persona guard replies ("Access Denied", "no") and the Qwen3 ⟨`think`⟩-block as incoherent. The reclassification recovers +12.6 pp of coherence the old classifier had hidden (ODILE aggregate 9.8% → 22.4%, n=500); the residual collapse is genuine degenerate over-jam, real and Qwen-concentrated (Llama-3.1-8B ODILE keeps 60% coherence, Qwen $\sim$ 9–21%). The split is therefore: Llama mild (−24 pp base→ODILE), Qwen severe (−55 to −69 pp), all on the adversarial canonical pre_prompts and all recovered by the neutral-pre_prompt fix. [♯]Qwen3-8B base coherence (77%, measured-only) corrects a double artifact in the original cross-model table (the reported 48/38 "coherent" were actually `B_leaked` counts from the ⟨`think`⟩ block restating the code); see the paper. [†]Llama-3.3-70B coherence is the original-classifier number. Canonical cell aggregator: `paper/latestresults/session_2026-05-20/cross_model/SUMMARY.md` §1c+; reclassification: `session_2026-05-26/tt_reclass/FINDINGS.md`. Visual summary in the paper Panel C.