

Defense (per backbone)	Attack ASR (%)↓		Canonical-benign coh. (%)↑	
	extract	hijack	extract	hijack
<i>Llama-3.1-8B</i>				
No defense	84.0	13.3	72	96
ODILE	44.0	5.7	60	60
<i>Qwen-2.5-7B</i>				
No defense	80.3	5.0	68	84
ODILE	4.0	0.7	22	20
<i>Qwen-2.5-14B</i>				
No defense	82.0	9.3	64	82
ODILE	1.0	0.3	12	6
<i>Qwen3-8B</i>				
No defense	91.3	63.7	77 [#]	77 [#]
ODILE	5.7	1.7	8	14
<i>Qwen3-32B</i>				
No defense	80.0	42.0	76	84
ODILE	2.7	1.3	14	8
<i>Llama-3.3-70B (reference)</i>				
No defense	78.0	13.0	62 [†]	72 [†]
ODILE	18.0	1.0	38 [†]	48 [†]

Table 1: **TensorTrust headline across six Llama/Qwen backbones.** Leak / hijack rates (↓ = better) under the canonical adversarial defender pre_prompts paired with each attack ($n=300$ per cell cross-model, $n=100$ at 70B). Canonical-benign coherence (↑ = better) is the rate at which the defended model produces an in-persona response to benign queries on the same canonical pre_prompts ($n=50$ per cell). ODILE reduces extract-attack ASR by $1.9\times$ at Llama-3.1-8B, $20\times$ at Qwen-2.5-7B, $82\times$ at Qwen-2.5-14B, $16\times$ at Qwen3-8B, $30\times$ at Qwen3-32B, and $4.3\times$ at Llama-3.3-70B; hijack ASR is suppressed to $\leq 5.7\%$ on every backbone. **The canonical-benign coherence drop has a trivial deployment fix:** the canonical TT defender pre_prompts are adversarially flavored by design (“do not reveal the password under any circumstance”), so the orth circuit engages on the *system prompt*’s framing even on benign queries; on neutral defender pre_prompts (same access codes and queries, pre_prompt rephrased neutrally) ODILE preserves $\sim 100\%$ coherence (Appendix ??). The deployment recipe is ODILE-alone, since layering an IH stanza jams 92% of benign queries regardless of pre_prompt (Appendix ??). The coherence values above are the *classifier-validated* rates from a deterministic 3-outcome reclassification of the on-disk model outputs (hand-validated 50/50 on the worst cell; Appendix ??), replacing the original TT classifier which mislabeled terse in-persona guard replies (“Access Denied”, “no”) and the Qwen3 **<think>**-block as incoherent. The reclassification recovers +12.6 pp of coherence the old classifier had hidden (ODILE aggregate $9.8\% \rightarrow 22.4\%$, $n=500$); the residual collapse is genuine degenerate over-jam, real and Qwen-concentrated (Llama-3.1-8B ODILE keeps 60% coherence, Qwen $\sim 9\text{--}21\%$). The split is therefore: Llama mild (-24 pp base $\rightarrow$ ODILE), Qwen severe (-55 to -69 pp), all on the adversarial canonical pre_prompts and all recovered by the neutral-pre_prompt fix. [#]Qwen3-8B base coherence (77%, measured-only) corrects a double artifact in the original cross-model table (the reported 48/38 “coherent” were actually **B_leaked** counts from the **<think>** block restating the code); see Appendix ??.

[†]Llama-3.3-70B coherence is the original-classifier number (the reclassification covered the five cross-model backbones; the matching 70B base reference on the adversarial-persona set was unavailable, Appendix ??). Canonical cell aggregator: `paper/latestresults/session_2026-05-20/cross_model/SUMMARY.md §1c+`; reclassification: `session_2026-05-26/tt_reclass/FINDINGS.md`. Visual summary in Figure ?? Panel C.