

Defense	dir	ign_p	imp_i	inj_a	f_mon	sys_m	sys_e	tk_d	tk_n	impl_s	sh_l	sh_e	sh_64	mean
No defense	8.86	23.24	29.57	17.15	26.51	11.16	10.99	16.90	20.00	4.14	3.51	8.75	1.69	14.04
ODILE	0.17	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.013
Meta-SecAlign-70B	1.80	2.73	2.26	2.97	1.04	0.87	2.28	2.97	3.38	1.39	2.15	1.74	1.80	2.11
Meta-SecAlign + Sandwich	2.15	2.84	2.67	2.95	1.74	1.56	2.56	2.50	3.56	2.32	3.08	2.08	1.74	2.44
Spotlight	10.02	11.38	27.16	8.56	20.70	6.40	15.47	10.60	12.23	9.46	8.02	8.88	6.55	11.96
Instruction Hierarchy	8.66	28.95	36.62	21.62	24.57	10.91	14.63	11.99	14.32	6.35	7.82	8.74	5.72	15.45
ODILE + Spotlight	—	—	—	—	—	—	—	—	—	—	—	—	—	0.20

Table 1: **Full AgentDojo injection-style matrix (Llama-3.3-70B)**. Average ASR (%) across the four AgentDojo suites for each standard attack family. The table includes the baseline, ODILE, Meta-SecAlign, prompt overlays, and prompt-stack comparators. ODILE reaches 0% ASR on 12 of 13 attack families; the residual is on `direct` (0.17%). Mean = 0.013%. MELON attack-family rows are in the paper. Column keys: `dir`=direct, `ign_p`=ignore_previous, `imp_i`=important_instructions, `inj_a`=injecagent, `f_mon`=fake_monologue, `sys_m`=system_message, `sys_e`=system_error_trap, `tk_d`=tk_descriptive, `tk_n`=tk_no_wrapper, `impl_s`=implicit_socratic, `sh_l`=shadowleak, `sh_e`=shadowleak_explicit, `sh_64`=shadowleak_b64. Per-cell denominators: banking $n=144$, slack $n=105$, workspace $n=240$, travel $n=140$ (4-suite mean $n=629$); Wilson 95% intervals and bootstrap protocol in the paper. The ODILE+Spotlight stacking-row aggregate is in the paper.