

Recipe variant	E-attacks ASR↓	Aug C/D ASR↓	InjecAgent ds_enh valid ASR↓
ODILE (augmentation-trained)	0.16	0.00	0.00
Meta-SecAlign-70B (reference baseline)	1.59	1.98	21.21

Table 1: **Augmentation-trained ODILE (Llama-3.3-70B headline)**. ODILE reaches 0.16% ASR on the AD-E grid ($n=3,145$), 0.00% on AD-Shift ($n=1,887$), and 0.00% ASR-valid on InjecAgent enhanced data-stealing ($n=1,054$). Structural argument: surface-form diversity at training time complements the intent-driven contrastive objective; a corpus dominated by a single wrapper style risks the adapter keying on the wrapper rather than the underlying intent.