

Attack family	Set	Cells	n	Base ASR/util	Meta-SecAlign ASR/util	ODILE ASR/util	
E aggregate	E	20	3,145	8.55 / 50.0		1.59 / 57.2	0.16 / 9.13
kernel_authority_mimicry	E	4	629	5.78 / 51.8		1.63 / 56.3	0.16 / 10.3
paraphrase_chain	E	4	629	7.02 / 50.3		1.91 / 56.6	0.16 / 11.0
refuse_to_refuse	E	4	629	20.07 / 47.6		1.04 / 57.0	0.16 / 10.3
tool_schema_redefinition	E	4	629	4.07 / 51.0		1.39 / 58.4	0.16 / 10.3
trickle_relay	E	4	629	5.83 / 49.4		1.97 / 57.8	0.16 / 10.6
Aug C/D aggregate	Aug	12	1,887	13.06 / 42.5		1.98 / 57.6	0.00 / 7.79
compositional_shadowmono	Aug	4	629	12.05 / 47.3		1.80 / 58.2	0.00 / 7.79
longctx_fakemono	Aug	4	629	20.51 / 36.1		1.98 / 58.8	0.00 / 7.79
longctx_shadowleak	Aug	4	629	6.61 / 44.1		2.15 / 55.7	0.00 / 7.79

Table 1: **Attack augmentations by family (Llama-3.3-70B)**. Each row reports ASR / utility-under-attack (%) over the four AgentDojo suites. The aggregate rows cover the five paraphrasing / authority-banner / refusal-coercion / tool-schema augmentations (E-attacks, $n=3,145$) and the three long-context augmentations (Aug, $n=1,887$). On the E-attacks, ODILE reaches 0.16% aggregate ASR; the per-attack residual concentrates on a single banking cell (`user_task_15` $\times$ `injection_task_4`) shared across multiple attack families, so ODILE’s behaviour is attack-content-invariant (§??). None of the eight augmentations appear in the training corpus (§??). Per-cell denominators match AgentDojo suite sizes (144/105/240/140). The utility column is under-attack utility (appendix view); benign utility is in Appendix ??.