Attack family	Set	Cells	n	Base ASR/util	Meta-SecAlign ASR/util	ODILE ASR/util
E aggregate	E	20	3,145	8.55 / 50.0	1.59 / 57.2	**0.16 / 9.13**
`kernel_authority_mimicry`	E	4	629	5.78 / 51.8	1.63 / 56.3	**0.16 / 10.3**
`paraphrase_chain`	E	4	629	7.02 / 50.3	1.91 / 56.6	**0.16 / 11.0**
`refuse_to_refuse`	E	4	629	20.07 / 47.6	1.04 / 57.0	**0.16 / 10.3**
`tool_schema_redefinition`	E	4	629	4.07 / 51.0	1.39 / 58.4	**0.16 / 10.3**
`trickle_relay`	E	4	629	5.83 / 49.4	1.97 / 57.8	**0.16 / 10.6**
Aug C/D aggregate	Aug	12	1,887	13.06 / 42.5	1.98 / 57.6	**0.00 / 7.79**
`compositional_shadowmono`	Aug	4	629	12.05 / 47.3	1.80 / 58.2	**0.00 / 7.79**
`longctx_fakemono`	Aug	4	629	20.51 / 36.1	1.98 / 58.8	**0.00 / 7.79**
`longctx_shadowleak`	Aug	4	629	6.61 / 44.1	2.15 / 55.7	**0.00 / 7.79**

Table 1: **Attack augmentations by family (Llama-3.3-70B).** Each row reports ASR / utility-under-attack (%) over the four AgentDojo suites. The aggregate rows cover the five paraphrasing / authority-banner / refusal-coercion / tool-schema augmentations (E-attacks, n=3,145) and the three long-context augmentations (Aug, n=1,887). On the E-attacks, ODILE reaches 0.16% aggregate ASR; the per-attack residual concentrates on a single banking cell (`user_task_15` × `injection_task_4`) shared across multiple attack families, so ODILE's behaviour is attack-content-invariant. None of the eight augmentations appear in the training corpus. Per-cell denominators match AgentDojo suite sizes (144/105/240/140). The utility column is under-attack utility (appendix view); benign utility is in the paper.