

Backbone (LoRA layers)	AD-std	AD-E	AD-Shift
	ASR / util	ASR / util	ASR / util
Llama-3.1-8B (L12–22)	8.7 / 17.5	3.3 / 19.2	6.8 / 17.2
+ ODILE	0.6 / 7.1	0.0 / 8–11	0.4 / 7.1
Qwen-2.5-7B (L10–19)	9.6 / 27.6	4.4 / 30.7	10.1 / 29.4
+ ODILE	0.0 / 11.7	0.0 / 8–11	0.0 / 11.7
Qwen3-8B (L13–25)	7.2 / 32.8	3.9 / 34.7	5.0 / 33.0
+ ODILE	0.0 / 8.4	0.0 / 8–11	0.0 / 8.4
Qwen-2.5-14B (L18–33)	8.3 / 44.7	3.0 / 46.1	3.2 / 46.4
+ ODILE	0.0 / 9.7	0.0 / 8–11	0.0 / 9.7
Qwen3-32B (L24–44)	8.1 / 43.2	4.2 / 38.0	5.5 / 35.9
+ ODILE	0.1 / 7.1	0.0 / 8–11	0.0 / 7.1

Table 1: **Cross-model transfer: 5 non-headline backbones $\times$ 3 AgentDojo benchmark families (under-attack-utility appendix view).** ASR / utility-under-attack (%). AD-std = 13×4 standard AgentDojo (52 cells); AD-E = 5×4 paraphrasing / authority-banner / refusal-coercion / tool-schema augmentations (20 cells); AD-Shift = 3×4 long-context augmentations (12 cells). None of the AD-E or AD-Shift attack families appear in the training corpus. ODILE drives ASR to $\leq 0.6\%$ on every measured cell. For AD-std pooled rates, each backbone has $n=8,177$ pairs, so near-zero ASR rows have tight Wilson intervals (e.g., a pooled 0.5% rate gives $\approx [0.37, 0.68]\%$). The cross-model claim is recipe-vs-no-defense consistency on each backbone; same-backbone head-to-head against external defenses (Meta-SecAlign, MELON, prompt overlays, CaMeL) is in the paper (70B only, since those comparators ship checkpoints for specific sizes). The under-attack utility numbers in this appendix view quantify residual non-malicious throughput once ODILE removes the attacker-target tool call from the agent’s output. The headline capability metric is benign-no-attack utility, tabulated in the paper.