

Backbone	Variant	Banking	Slack	Travel	Workspace	Overall
		ASR / util	ASR / util	ASR / util	ASR / util	ASR / util
Llama-3.1-8B	no defense	14.9 / 21.9	19.7 / 29.6	1.0 / 1.9	0.1 / 17.6	8.7 / 17.5
	ODILE	2.4 / 16.1	0.0 / 9.6	0.0 / 0.1	0.0 / 2.4	0.6 / 7.1
Qwen-2.5-7B	no defense	11.2 / 40.3	19.7 / 29.5	5.0 / 8.6	2.5 / 32.0	9.6 / 27.6
	ODILE	0.0 / 25.0	0.0 / 16.7	0.0 / 0.0	0.0 / 5.0	0.0 / 11.7
Qwen3-8B	no defense	5.2 / 43.9	17.8 / 23.9	5.3 / 26.7	0.7 / 36.7	7.2 / 32.8
	ODILE	0.0 / 24.7	0.0 / 4.8	0.0 / 1.2	0.0 / 2.9	0.0 / 8.4
Qwen-2.5-14B	no defense	9.6 / 58.4	19.3 / 42.0	3.1 / 25.8	1.4 / 52.5	8.3 / 44.7
	ODILE	0.0 / 25.0	0.0 / 15.2	0.0 / 0.6	0.0 / 2.5	0.0 / 9.7
Qwen3-32B	no defense	8.5 / 50.2	18.8 / 43.8	4.8 / 27.6	0.3 / 51.4	8.1 / 43.2
	ODILE	0.2 / 19.0	0.0 / 6.7	0.0 / 0.0	0.0 / 2.7	0.1 / 7.1

Table 1: **AgentDojo-standard per-suite breakdown for the five cross-model backbones (under-attack utility, appendix view)**. ASR / utility-under-attack (%); each suite cell averages 13 attack families on the AD-standard grid. Overall is the 52-cell mean. ODILE drives ASR to $\leq 0.6\%$ on every (backbone, suite) pair. Under-attack utility is the residual non-malicious throughput once ODILE removes the attacker-target tool call from the agent’s output (the adapter outputs jam under attack; §??). The benign-no-attack headline metric is in Table ??. Per-suite denominators (144/105/240/140); Wilson and bootstrap protocol in Appendix ??.