

Defense (Llama-3.1-8B)	ASR (%)↓	Benign util. (%)↑	Inference cost
No defense (base)	11.1	25.0	1×
PromptGuard detector [ref]	11.8	25.0	1× + classifier
PromptGuard detector ($t=0.5$)	10.4	25.0	1× + classifier
Sanitizer (LLM-rewriter) [ref]	2.1	18.8 (−6.2)	2× (rewrite/tool out.)
ODILE	2.1	25.0 (−0.0)	1× (LoRA)

Table 1: **Static firewall comparators vs. ODILE at Llama-3.1-8B** (AgentDojo banking, `important_instructions`, $n=144$ attack / $n=16$ benign). The detector firewall (PromptGuard) does not change ASR meaningfully — it scores the polite `important_instructions` frame ~ 0.23 vs. ~ 0.999 on blatant overrides, so recall on stylistically-disguised agentic injections is the failure mode [ref]. The Sanitizer matches ODILE on ASR (2.1%) but costs 6.2pp benign utility (it over-sanitizes legitimate tool outputs) and an extra LLM call per tool output; ODILE matches the same ASR at zero benign-utility loss and 1× cost. Static parity breaks under *adaptive* attack: the Sanitizer is cracked 15/48 across two backbones while ODILE holds 0/64 (Table ??). The *Sanitizer* is our faithful re-implementation of Bhagwatkar et al. [ref] using a local Llama-3.1-8B rewriter (their original 0% AgentDojo number uses GPT-4o / 70B-class rewriters); workspace-suite static cells in Appendix ??.