

	<u>dh_base</u>		<u>dh_enh.</u>		<u>ds_base</u>		<u>ds_enh.</u>	
Defense	all	valid	all	valid	all	valid	all	valid
No defense	47.06	55.68	84.71	87.80	70.87	79.09	96.75	98.64
Meta-SecAlign-70B	0.20	3.57	1.18	24.00	0.55	5.36	2.59	21.21
ODILE	0.20	2.17	0.00	0.00	0.00	0.00	0.00	0.00
ODILE + fav	0.00	0.00	0.00	0.00	—	—	0.00	0.00
ODILE + fav + def	0.00	0.00	0.00	0.00	0.00	0.00	—	—

Table 15: **InjecAgent full table, Llama-3.3-70B**. Each cell reports ASR-all / ASR-valid (%) over $n=1,054$ test cases per setting. ASR-valid = successes / parseable outputs; the meaningful robustness column when invalid-rates differ. ODILE reaches $\leq 0.20\%$ ASR-all on every setting. With inference-time stanzas (+fav, +def) the residual **dh_base** cell drops to 0%. Cross-model InjecAgent in Table ??; benign InjecAgent at 70B for ODILE-alone: 88% correct on $n=100$ legitimate tool-call items.