Attack	slack			workspace		
	ASR↓	Util↑	Refusal	ASR↓	Util↑	Refusal
compositional_shadowmono	**0.133**	0.486	0.14	0.000	0.438	0.03
kernel_authority_mimicry	0.019	0.438	0.43	0.000	0.471	0.03
longctx_fakemono	0.000	0.314	**0.65**	0.008	0.392	0.02
paraphrase_chain	0.029	0.352	0.33	0.000	0.463	0.02
refuse_to_refuse	0.019	0.286	0.57	0.008	0.467	0.03
tool_schema_redefinition	0.019	0.438	0.35	0.004	0.446	0.03
trickle_relay	0.029	0.524	0.13	0.000	0.458	0.01

Table 17: **MELON × AD-E + AD-Shift attacks (Llama-3.3-70B base, 14 cells).** ASR, under-attack utility, and MELON-induced refusal rate per (attack, suite) cell. Workspace ASR is ≤ 0.008 on every cell; slack carries every measurable failure. The single highest cell is `compositional_shadowmono` × `slack` at 13.3% ASR, where masked-prompt contrastive detection sees similar tool-call sequences with and without the injection. Slack-suite refusal rates (13–65%) compound the structural over-refusal already present on benign slack data (31% MELON-induced refusal on slack, see §**??**). Per-cell denominator n=144 banking, n=105 slack, n=240 workspace, n=140 travel; with only 14 (attack, suite) cells covered here and several at point-estimate 0, single-cell Wilson intervals are wide (e.g., the 13.3% slack cell has a 95% Wilson interval of $\approx [7.9, 21.5]\%$). Treat individual MELON × attack-family cells as directional, not significance-tested. Full protocol in Appendix **??**.