

Attack family	Base	Base+M	ODILE	ODILE+M	Meta-SecAlign	Meta-SecAlign+M
fake_monologue	26.51 / 42.4	0.00 / 25.7	0.00 / 8.3	0.17 / 12.4	1.04 / 67.6	0.52 / 36.9
important_instructions	29.57 / 46.7	0.53 / 25.5	0.00 / 8.3	0.17 / 11.7	2.26 / 67.2	0.41 / 34.5
shadowleak	3.51 / 45.5	0.41 / 28.8	0.00 / 8.3	0.17 / 11.3	2.15 / 65.0	0.35 / 34.5
tk_no_wrapper	20.00 / 43.3	5.4 / 33.0	0.00 / 8.3	0.17 / 11.3	3.38 / 66.9	0.93 / 36.9
Mean	19.90 / 44.5	0.24 / 28.0	0.00 / 8.3	0.17 / 11.7	2.21 / 66.7	0.55 / 35.7

Table 1: **Per-attack MELON breakdown (Llama-3.3-70B, appendix view)**. Entries are cell-average ASR / utility-under-attack (%) across the four AgentDojo suites for each attack family. “+M” denotes MELON layered on the same defense. The bottom **Mean** row is the per-defense aggregate. ODILE reaches 0% ASR on every attack family. Meta-SecAlign-70B occupies a different region of the same frontier: higher under-attack utility (66.7% mean) at higher ASR (2.21% mean). Layering MELON on ODILE does not improve ASR over ODILE-alone. Each (attack family, defense) cell is the mean over four AgentDojo suites ($n=629$ pairs / cell); Wilson and bootstrap protocol in the paper.