

Stack	ASR ↓	Util ↑	Δ ASR	Δ util	Read
ODILE	0.01	8.3	—	—	jam output under attack
ODILE + Spotlight	0.20	25.0	+0.2pp	+16.7pp	same-cell utility gain
ODILE + Instruction Hierarchy	0.69	20.4	+0.7pp	+12.1pp	same-cell ASR cost; util gain
ODILE + Repeat	2.08	21.8	+2.1pp	+13.5pp	same-cell ASR cost; util gain
Meta-SecAlign-70B	2.11	66.6	—	—	full-DPO comparator
Meta-SecAlign + Sandwich	2.44	70.0	+0.3pp	+3.4pp	full-coverage overlay stack
Meta-SecAlign + Instruction Hierarchy	4.97	69.7	+2.9pp	+3.1pp	same-cell ASR cost
Meta-SecAlign + Spotlight	5.40	62.6	+3.3pp	−4.0pp	same-cell utility cost
No defense: Llama-3.3-70B-Instruct	14.04	43.9	—	—	no defense
Spotlight	11.96	48.5	−2.1pp	+4.6pp	full-coverage overlay on base
Instruction Hierarchy	15.45	50.7	+1.4pp	+6.7pp	full-coverage overlay on base
Repeat-Prompt	13.89	51.4	−0.2pp	+6.3pp	1-cell sample; coverage-limited
DeBERTa classifier	4.59	19.7	—	—	classifier; ASR drop with util collapse

Table 19: **Overlay-defense stacking matrix on AgentDojo standard (Llama-3.3-70B; under-attack-util appendix view)**. Average ASR and utility-under-attack (%). Δ columns are same-covered-cell changes against the corresponding unstacked defense. Stacking overlays on ODILE raise under-attack utility by +12 to +17pp while leaving ASR at $\leq 2.1\%$; the deployment-recipe choice across these rows depends on the application’s ASR threshold. Coverage is full 52 cells unless flagged in the “Read” column. The corresponding benign-utility numbers under each stack (the headline capability metric) are in Appendix ??.