

Benchmark	Llama-3.1-8B	Qwen-2.5-7B	Qwen-2.5-14B	Qwen3-8B	Qwen3-32B	Llama-3.3-70B
AlpacaEval LC ($n=805$)	32.0 $\rightarrow$ 31.6 -0.4	38.7 $\rightarrow$ 38.2 -0.5	44.8 $\rightarrow$ 35.5 -9.3	19.1 $\rightarrow$ 21.1 $+2.1$	79.8 $\rightarrow$ 61.4 -18.5	48.6 $\rightarrow$ 47.9 -0.7
GSM8K ($n=200$)	85.5 $\rightarrow$ 83.5 -2.0	89.0 $\rightarrow$ 87.5 -1.5	91.5 $\rightarrow$ 90.5 -1.0	95.5 $\rightarrow$ 94.5 -1.0	97.5 $\rightarrow$ 71.5 -26.0	95.0 $\rightarrow$ 95.5 $+0.5$
IFEval strict ($n=200$)	62.5 $\rightarrow$ 59.5 -3.0	54.5 $\rightarrow$ 51.5 -3.0	64.0 $\rightarrow$ 52.0 -12.0	64.5 $\rightarrow$ 57.0 -7.5	66.5 $\rightarrow$ 27.0 -39.5	73.5 $\rightarrow$ 71.0 -2.5
BFCL non-live AST ($n=1240$)	73.98 $\rightarrow$ 73.35 -0.63	74.98 $\rightarrow$ 75.12 $+0.14$	—	—	—	—
τ -bench airline ($n=50$)	42.0 $\rightarrow$ 38.8 -3.2	—	—	—	—	—

Table 1: **General-purpose capability retention under ODILE.** Each cell is base $\rightarrow$ **ODILE** with the signed change (**green** $+$ = **gain**, **red** $-$ = **drop**). Capability is preserved within a few points at 8B (Llama-3.1-8B, Qwen-2.5-7B, Qwen3-8B) and 70B-Llama; Qwen-2.5-14B and especially Qwen3-32B carry a real instruction-format cost (the Qwen3-32B cells are corrected no-think re-runs; thinking-mode measurement artifacts are documented in the paper). Qwen3-8B AlpacaEval LC and Qwen-2.5-7B BFCL *improve*. BFCL/ τ -bench on the remaining backbones is a camera-ready commitment. Full per-subset numbers in the paper.