Defense (Llama-3.1-8B)	AD-std ASR / util	AgentDyn benign	IA-enh ASR / valid	TT-extract leak	WASP K8 strict	BFCL retain
base	16% / 26%	8.3%	48.0% / 61.0%	10%	$\sim$ 14%	100% (ref)
ODILE	**0.9**% / **29**%	0%‡	0.0% / **43.7**%	14%	0%	93% (−4.5)
ReasAlign	0.4% / 13%	1.7%‡	0.0% / 0.3%†	4%	0%	93% (−4.1)

Cross-benchmark model-level-defense comparator at Llama-3.1-8B. ODILE vs. base and ReasAlign across all five benchmarks. ReasAlign matches ODILE on ASR but costs $\sim$ 13 pp AD-standard benign utility and $\sim$ 4 pp BFCL; ODILE preserves both benign utility and, critically, InjecAgent parseability (valid-rate 43.7% vs. ReasAlign's 0.3%). The 70B headline comparator is Meta-SecAlign (the 70B Pareto table); ReasAlign is the 8B representation-class comparator here.

† valid-rate ≈ 0 means the defense *over-jams* InjecAgent: ASR is trivially 0 because no parseable output survives. ODILE is the only arm that kills InjecAgent ASR while *preserving* parseability (valid-rate 43.7% vs. base 61.0%), i.e. a genuine defense rather than defense-by-output-mangling. ‡ AgentDyn benign at Llama-3.1-8B is the known over-jam regression (the λ_4 harmful-looking-benign anchor fix exists for Qwen3-8B and is not yet transferred to 8B).