

Defense (Llama-3.1-8B)	AD-std	AgentDyn	IA-enh	TT-extract	WASP K8	BFCL
	ASR / util	benign	ASR / valid	leak	strict	retain
base	16% / 26%	8.3%	48.0% / 61.0%	10%	$\sim 14\%$	100% (ref)
ODILE (nullify, headline)	0.9% / 29%	0% [‡]	0.0% / 43.7%	14%	0%	93% (−4.5)
ODILE (triplet, ablation)	0.2% / 32%	3.3% [‡]	0.0% / 0.1% [†]	8%	0%	101%
ReasAlign	0.4% / 13%	1.7% [‡]	0.4% / 0.3% [†]	4%	0%	93% (−4.1)

Cross-benchmark model-level-defense comparator at Llama-3.1-8B. ODILE (nullify headline + triplet ablation) vs. base and ReasAlign across all five benchmarks. ReasAlign matches ODILE on ASR but costs ~ 13 pp AD-standard benign utility and ~ 4 pp BFCL; ODILE-triplet occupies the Pareto corner (lowest AD-std ASR at highest benign utility, lowest TT-extract leak, capability-preserving on BFCL). The 70B headline comparator is Meta-SecAlign (the 70B Pareto table); ReasAlign is the 8B representation-class comparator here.

[†] valid-rate ≈ 0 means the defense *over-jams* InjecAgent: ASR is trivially 0 because no parseable output survives. ODILE-nullify is the only arm that kills InjecAgent ASR while *preserving* parseability (valid-rate 43.7% vs. base 61.0%), i.e. a genuine defense rather than defense-by-output-mangling. [‡] AgentDyn benign at Llama-3.1-8B is the known over-jam regression (the λ_4 harmful-looking-benign anchor fix exists for Qwen3-8B and is not yet transferred to 8B).