

Defense (Llama-3.1-8B)	AgentDojo ASR %↓	Benign utility %↑
No defense	16.0	26.0
ODILE	0.9 −15.1	29.0 +3.0
ReasAlign	0.4 −15.6	13.0 −13.0
Meta-SecAlign-70B	—	—
Meta-SecAlign + Sandwich	—	—
MELON detector	—	—
Instruction Hierarchy	—	—
Spotlight	—	—
Repeat-Prompt	—	—
DeBERTa classifier	—	—
PromptGuard [†]	—	—
Sanitizer [†]	—	—
CaMeL	—	—
<i>Broader published field (cross-suite means, provisional):</i>		
Spotlight-encode	0.2 −15.8	15.0 −11.0
Spotlight-datamark	3.9 −12.1	19.0 −7.0
PromptArmor	0.2 −15.8	19.0 −7.0
TaskShield	0.2 −15.8	23.8 −2.2
Perplexity filter	1.0 −15.0	12.7 −13.3

Table 1: **AgentDojo** — **Llama-3.1-8B** (comparator-sweep grid; *provisional*). Badge = change vs. no-defense. ODILE holds the highest benign utility (29%) among the low-ASR defenses while matching them on ASR. [†]PromptGuard/Sanitizer were measured on a banking subset (different grid, base 11.1%); see the firewall table. “—” = comparator not run at this scale (70B-only at submission; cross-backbone is camera-ready).