

Defense (Llama-3.1-8B)	AgentDojo ASR %↓	Benign utility %↑
No defense	16.0	26.0
ODILE	0.9 ↓ _{15.1}	29.0 ↑ _{3.0}
ReasAlign	0.4 ↓ _{15.6}	13.0 ↓ _{13.0}
Meta-SecAlign-70B	—	—
Meta-SecAlign + Sandwich	—	—
MELON detector	—	—
Instruction Hierarchy	—	—
Spotlight	—	—
Repeat-Prompt	—	—
DeBERTa classifier	—	—
PromptGuard [†]	—	—
Sanitizer [†]	—	—
CaMeL	—	—
<i>Broader published field (cross-suite means, provisional):</i>		
Spotlight-encode	0.2 ↓ _{15.8}	15.0 ↓ _{11.0}
Spotlight-datamark	3.9 ↓ _{12.1}	19.0 ↓ _{7.0}
PromptArmor	0.2 ↓ _{15.8}	19.0 ↓ _{7.0}
TaskShield	0.2 ↓ _{15.8}	23.8 ↓ _{2.2}
Perplexity filter	1.0 ↓ _{15.0}	12.7 ↓ _{13.3}

Table 2: **AgentDojo** — **Llama-3.1-8B** (comparator-sweep grid; *provisional*). Badge = change vs. no-defense. ODILE holds the highest benign utility (29%) among the low-ASR defenses while matching them on ASR. [†]PromptGuard/Sanitizer were measured on a banking subset (different grid, base 11.1%); see the firewall table. “—” = comparator not run at this scale (70B-only at submission; cross-backbone is camera-ready).