

Defense (Llama-3.3-70B)	AgentDojo ASR %↓	Benign utility %↑
No defense	14.04	59.9
ODILE	0.01 ↓ _{14.0}	51.5 ↓ _{8.4}
Meta-SecAlign-70B	2.11 ↓ _{11.9}	62.6 ↑ _{2.7}
Meta-SecAlign + Sandwich	2.44 ↓ _{11.6}	63.4 ↑ _{3.5}
MELON detector	0.24 ↓ _{13.8}	48.6 ↓ _{11.3}
Instruction Hierarchy	15.45 ↑ _{1.4}	59.4 ↓ _{0.5}
Spotlight	11.96 ↓ _{2.1}	60.0 ↑ _{0.1}
Repeat-Prompt	13.89 ↓ _{0.2}	51.4 ↓ _{8.5}
DeBERTa classifier	4.59 ↓ _{9.5}	59.2 ↓ _{0.7}
PromptGuard	—	—
Sanitizer	—	—
ReasAlign	—	—
CaMeL	<i>partial: 2/4 suites, 0% ASR / 0.50–0.90 benign</i>	

Table 1: **AgentDojo headline — Llama-3.3-70B.** Large black number = absolute value; small badge (bottom-right) = change vs. no-defense (**green**=improvement / **red**=regression; ↓/↑=direction; for ASR a drop is the improvement). ODILE reaches the lowest ASR (0.01%) at 1× inference cost; Meta-SecAlign holds higher benign utility at ∼200× the ASR. $n=8,177$ pairs (ASR), $n=97$ (benign). Blank = comparator not run at this scale.