

Defense (Llama-3.3-70B)	AgentDojo ASR %↓	Benign utility %↑
No defense	14.04	59.9
ODILE	0.01 −14.0	51.5 −8.4
Meta-SecAlign-70B	2.11 −11.9	62.6 +2.7
Meta-SecAlign + Sandwich	2.44 −11.6	63.4 +3.5
MELON detector	0.24 −13.8	48.6 −11.3
Instruction Hierarchy	15.45 +1.4	59.4 −0.5
Spotlight	11.96 −2.1	60.0 +0.1
Repeat-Prompt	13.89 −0.2	51.4 −8.5
DeBERTa classifier	4.59 −9.5	59.2 −0.7
PromptGuard	—	—
Sanitizer	—	—
ReasAlign	—	—
CaMeL	<i>partial: 2/4 suites, 0% ASR / 0.50–0.90 benign</i>	

Table 1: **AgentDojo headline — Llama-3.3-70B.** Large black number = absolute value; small badge (bottom-right) = change vs. no-defense (**green**=improvement / **red**=regression; ↓/↑=direction; for ASR a drop is the improvement). ODILE reaches the lowest ASR (0.01%) at 1× inference cost; Meta-SecAlign holds higher benign utility at ∼200× the ASR. $n=8,177$ pairs (ASR), $n=97$ (benign). Blank = comparator not run at this scale.