

Defense (Qwen3-32B)	AgentDojo ASR %↓	Benign utility %↑
No defense	8.1	35.3
ODILE	0.5 ↓7.6	37.5 ↑2.2
Meta-SecAlign-70B	—	—
Meta-SecAlign + Sandwich	—	—
MELON detector	—	—
Instruction Hierarchy	—	—
Spotlight	—	—
Repeat-Prompt	—	—
DeBERTa classifier	—	—
PromptGuard	—	—
Sanitizer	—	—
ReasAlign	—	—
CaMeL	—	—

Table 6: **AgentDojo** — **Qwen3-32B** (52-cell, $n=8,177$ ASR / $n=629$ benign). ODILE drives ASR to 0.5% *and* raises benign agentic utility by +2.2 pp (the parser-strict capability drops are a separate thinking-mode story, not the agentic-benign axis). “—” = comparator not run at this scale.