

Defense (Qwen-2.5-14B)	AgentDojo ASR %↓	Benign utility %↑
No defense	8.3	46.1
ODILE	0.0 −8.3	41.8 −4.3
CaMeL [‡]	0.0 [‡]	0.0 [‡]
Meta-SecAlign-70B	—	—
Meta-SecAlign + Sandwich	—	—
MELON detector	—	—
Instruction Hierarchy	—	—
Spotlight	—	—
Repeat-Prompt	—	—
DeBERTa classifier	—	—
PromptGuard	—	—
Sanitizer	—	—
ReasAlign	—	—

Table 1: **AgentDojo** — **Qwen-2.5-14B** (52-cell, $n=8,177$ ASR / $n=629$ benign). [‡]CaMeL’s 0% ASR / 0% benign is *structurally vacuous* (scale-floor: the privileged LLM cannot drive the restricted-Python interpreter, so no tool calls are produced either direction) — not an earned defense point. ODILE is the only defense delivering *useful* zero-ASR at this scale. “—” = comparator not run at this scale.