

Table 8: **WASP (web-agent prompt injection): verified-tight adaptive attacker-action rate / benign utility (%)**, by backbone. WASP [ref] is an agentic web prompt-injection benchmark; we run a faithful proxy of it — the 21 canonical attacker-goal records and injection-format templates are preserved verbatim, with a 5-tool stub standing in for the live Dockerized web environment (canonical containerized WASP is a camera-ready commitment; Appendix ??). Attackers: K=8 paraphrase + TAP3×8 iterative (up to 24 attempts/record); success = a record-specific attacker fingerprint in a non-browse tool-call’s arguments. “—” = not run (Sanitizer not run at 70B). Aggregate is crack-count/attempts (ODILE denominator includes both 70B recipe variants). Qwen3-8B: static-only, ODILE 0/40 (adaptive not run).

backbone	base ASR / ben	+Sanitizer ASR / ben	+ODILE ASR / ben
Llama-3.1-8B	0 / 100	0 / 100	0 / 83
Qwen-2.5-7B	35 / 100	15 / 98	0 / 95
Qwen-2.5-14B	15 / 69	5 / —	0 / 69
Qwen3-32B	7.5 / 19	0 / 12	0 / 33
Llama-3.3-70B	5 / 100	—	0 / 93 (n), 0 / 100 (t)
aggregate ASR	25/200 (12.5%)	8/160 (5.0%)	0/239 (0%)