

Recipe (Qwen3-8B)	AD-std ASR % $\downarrow$	Benign util % $\uparrow$	train rows
ODILE (exclude 3 pairs)	0.12	40.32	844
ODILE (include 3 pairs)	0.29 _{+0.17}	39.95 _{−0.37}	850

Table 1: **Worsened-pairs include-vs-exclude ablation.** The only difference between rows is the three structurally-ambiguous banking/user_task_15 pairs (attacker IBAN = legitimate recipient, so the harmful and benign twins are byte-identical). Including them is pure downside: AD-standard ASR rises (0.12 $\rightarrow$ 0.29%, +0.17pp; direct + ignore_previous + important_instructions across four suites) while benign utility is essentially unchanged (−0.37pp). Every worsened training sample traces to this one task, which is why we exclude it by a structural rule rather than a per-sweep statistic. Big number = value; small colored delta = change vs. the excluded (ODILE) recipe.